

Group 14 Progress Report: Municipal Waste Image Classifier

Sunny Yao, Aswin Kuganesan, Yousef Shahin

{yaoh24, kugana7, shahiy3}@mcmaster.ca

1 Introduction

The rapid increase in municipal solid waste has become a major global issue, contributing to landfill overflow, environmental pollution, and the depletion of natural resources. Effective waste management plays a vital role in addressing these challenges by promoting recycling, conserving resources, and reducing the overall environmental footprint. However, traditional methods of waste sorting are often manual, time-consuming, and error-prone, limiting the efficiency and scalability of recycling systems.

This report presents the development of a Municipal Waste Image Classification system that applies deep learning techniques to automate the identification and categorization of waste materials. The system is designed to classify common waste types such as plastic, paper, glass, metal, and organic matter based on image data. By automating the classification process, the project aims to improve recycling efficiency, streamline waste management operations, and support environmental sustainability efforts.

The main challenge of this project lies in the variability and complexity of real-world waste images. Factors such as lighting conditions, background noise, and material degradation make accurate classification difficult. Overcoming these challenges requires robust machine learning models capable of handling diverse input data while maintaining high accuracy and efficiency. The success of this project would demonstrate the potential of artificial intelligence to transform waste management and support more sustainable, environmentally responsible cities.

Image classification for waste management has gained significant attention in recent years, with various studies leveraging deep learning architectures to improve recycling efficiency. Convolutional Neural Networks (CNNs) have proven espe-

cially effective for this task due to their ability to automatically learn hierarchical feature representations from raw image data.

Several studies have demonstrated the potential of CNNs for automated waste sorting. ? developed a deep learning-based waste classification system using transfer learning with pre-trained models such as VGG16 and ResNet50, achieving accuracies exceeding 90% on a multi-class dataset. Their work showed that fine-tuning deep architectures on domain-specific waste imagery can substantially enhance performance across diverse material types including plastic, paper, and metal.

Earlier, ? constructed one of the first comprehensive waste image datasets, known as TrashNet, and implemented a custom CNN model for recyclability classification. They identified key challenges in differentiating visually similar materials, such as glass and plastic, and emphasized the importance of data augmentation and preprocessing for improving model robustness.

Building on this foundation, ? proposed an ensemble CNN approach that combines the outputs of multiple architectures to improve overall accuracy and generalization. Their ensemble method effectively mitigated misclassification between visually overlapping categories by leveraging complementary feature extraction from different models.

To enhance computational efficiency, particularly for real-time applications, ? introduced a lightweight CNN model optimized for deployment on edge devices. Their framework balanced model complexity with inference speed, enabling efficient on-device waste classification without requiring high-end hardware.

More recently, ? integrated deep learning with Internet of Things (IoT) technologies to develop an intelligent waste classification system for smart cities. Their model demonstrated the feasibility of combining AI-driven image recognition with connected infrastructure to support sustainable urban

waste management.

Our approach builds upon these foundational works by implementing a custom CNN architecture that progressively extracts features through multiple convolutional blocks. It incorporates global average pooling to minimize overfitting and dropout for regularization. While our dataset is currently limited in size, our model design follows proven principles from the literature and is optimized for six-category municipal waste classification.

2 Dataset

We used the garbage classification dataset from Kaggle and downloaded it to set up the dataset. This dataset was chosen because it had a large number of images and had clearly separated folders for each class (garbage category), which meant that the data did not need to be manually annotated. Furthermore, this meant that there was no human error since the annotation was not done manually by us. Since there was a large amount of data for each category, this meant that there was equal representation for each class. Therefore, we did not need to do any preprocessing or adjust the model to account for any class imbalances. The dataset is organized into 6 classes: cardboard, metal, paper, glass, plastic, and trash.

We created a `GarbageDataset` class that first looks inside the dataset directory and finds all the folders which are named after the 6 classes. Then, it maps the classes from numbers 0 to 5 so that each class can be represented as a number. It traverses through each of the folders and traverses every image to check if they are in JPG, PNG or JPEG format. If the image is in the correct format, it will add the image to a list of images. This list contains all the image paths and the labels of all of the images. This list is used by the model as the dataset, since it contains the image path, so it can access the image, and its corresponding label, so the model knows what class the image belongs to. This class can return the number of images in the dataset through `__len__`. It can also provide any image in the dataset with its corresponding label through `__getitem__` if given an index as a parameter. The image gets converted to RGB, which ensures that all images have three colour channels, even if the original images are in a different format, which is important because the Convolutional Neural Network expects a constant number of channels. If the channel dimensions are inconsistent,

the training will be prone to error. Furthermore, the image can be transformed (resizing, normalization, or augmentation) if the `GarbageDataset` class is instantiated with a transform defined.

The preprocessing operations performed on the dataset are done using PyTorch's transforms module. The training images are resized randomly and cropped to 224x224 pixels, and the images were augmented through random horizontal flipping to improve generalization. The images were then converted to tensors for the model to process them properly. Finally, the images are normalized using the mean and standard deviation of a pretrained model to stabilize gradient updates and reduce variation in pixel range through standardization. The training images also undergo image processing by cropping the images to 224x224 pixels. Furthermore, the images are converted to tensors and normalized.

Our previous model only used thirty percent of our dataset, which led to a lower accuracy during training. The model is expected to be more accurate with more data available, since it can determine more patterns and similarities between images of the same class if it has more data to train on. Therefore, we used the entire dataset for this model, which is expected to increase the accuracy and runtime of our model per epoch. An 80 percent training and 20 percent testing split was used to evaluate the model, and stratification is applied to ensure that the same class distribution is used for both the training and testing datasets. Finally, data loaders are created for the training and testing datasets to allow for efficient image processing and loading, with batch sizes of 32 and with the training dataset being randomized for each epoch. This improved learning robustness since it ensured that the model was learning after each epoch and that the accuracy of the model was not dependent on the training dataset.

3 Features

The primary inputs to our models are RGB images of municipal solid waste from the Kaggle "Garbage Dataset Classification" collection. Each image belongs to exactly one of six categories (cardboard, glass, metal, paper, plastic, or trash), and the folder structure encodes the class labels. We implemented a custom PyTorch Dataset class (`GarbageDataset`) that recursively traverses the dataset root directory, identifies all

valid image files (PNG, JPG, JPEG), and constructs a list of $(image_path, label_index)$ pairs. The `label_index` is an integer derived from a class-to-index mapping built from the folder names. All images are loaded via PIL and converted to RGB to enforce a consistent three-channel input, which is necessary because the convolutional neural network (CNN) expects a fixed number of channels; this also avoids runtime errors due to grayscale or otherwise inconsistent image formats.

For each image, we perform a series of feature preprocessing and representation-learning-oriented steps using `torchvision.transforms`. The raw height and width of the images vary across the dataset, so they are first standardized to a resolution compatible with our CNN. In the final configuration, we use a training transform (`train_tf`) consisting of `RandomResizedCrop(224, scale=(0.8, 1.0))`, `RandomHorizontalFlip()`, `ToTensor()`, and `Normalize(mean, std)`. The random resized crop changes both the scale and crop location of the image while always producing a 224×224 output, which encourages the model to learn features that are robust to object size and framing. The random horizontal flip introduces left-right invariances, which are appropriate because flipping an image of a bottle or a can does not change its semantic class. `ToTensor()` converts the PIL image to a tensor in channel–height–width format and scales pixel intensities to the $[0, 1]$ range. Finally, we apply channel-wise normalization using ImageNet statistics ($\text{mean} = [0.485, 0.456, 0.406]$, $\text{std} = [0.229, 0.224, 0.225]$), which centers and scales each channel to stabilize optimization and leverage best practices from large-scale image classification.

For evaluation, we adopt a deterministic test transform (`test_tf`) composed of `Resize(256)`, `CenterCrop(224)`, `ToTensor()`, and the same `Normalize(mean, std)`. This pipeline ensures that all test images are processed consistently, without randomness, so that performance metrics are comparable across runs. The resize and center crop standardize the spatial dimensions while preserving as much of the object as possible, and the normalization aligns the distribution of test inputs with that of the training data.

We did not apply explicit hand-crafted feature engineering (e.g., SIFT, HOG, or color histograms); instead, we rely on the CNN to perform representation learning directly from pixel intensities. The convolutional layers automatically learn hier-

archies of features (edges, textures, shapes, and higher-level object parts) tailored to the waste classification task. However, we did experiment with different feature preprocessing and augmentation schemes. Early experiments used only simple `Resize(224, 224)` and `RandomHorizontalFlip` as the training transform, with minimal augmentation. This configuration produced reasonable performance but exhibited signs of overfitting as training accuracy increased faster than validation accuracy. We then introduced stronger spatial augmentation via `RandomResizedCrop`, and in some exploratory runs added color jitter and small random rotations to encourage robustness to lighting changes and orientation. These more aggressive augmentations sometimes reduced overfitting but could also make training more challenging when combined with strong regularization (dropout and weight decay).

We did not perform traditional feature selection in the tabular sense, because the model operates directly on images. Instead, our “feature selection” and “feature variation” experiments were realized through changes to the input pipeline and model capacity. Concretely, we varied (i) the proportion of the dataset used (e.g., training on 30% vs. 100% of the images), (ii) the intensity of data augmentation (light vs. heavy transforms), and (iii) the effective depth of the CNN (using three convolutional blocks vs. four). These variations allowed us to study how different input representations and learned feature hierarchies affect generalization performance. Our best-performing configuration, which we use as the final model, employs three convolutional blocks, moderate augmentation (random resized crop and horizontal flip), and ImageNet-style normalization, yielding a validation accuracy of approximately 84.75%.

4 Implementation

The core model used in this project is a custom convolutional neural network implemented in PyTorch, named `WasteClassifierModelV1`. The network architecture follows a conventional deep vision design: a stack of convolutional blocks for hierarchical feature extraction, followed by global average pooling and a fully connected classifier. The model accepts input tensors of shape $(batch_size, 3, 224, 224)$ and outputs logits for C waste classes, where C is the number of distinct garbage categories discovered from the dataset

folder structure.

In its final configuration, `WasteClassifierModelV1` consists of three convolutional blocks. The first block applies two 3×3 convolutions with `hidden_units` output channels (set to 64 in our experiments), each followed by a ReLU nonlinearity, and then a 2×2 max-pooling layer to reduce spatial resolution by a factor of two. The second block repeats this pattern with doubled channel depth (from 64 to 128), again followed by max pooling. The third block applies a single 3×3 convolution increasing the channels from 128 to 256, followed by ReLU and 2×2 max pooling. A fourth convolutional block is defined in the code, which would further increase the channels from 256 to 512; however, we found that including this block, especially in combination with strong regularization, did not improve validation performance on the available data, and in some runs actually degraded generalization. As a result, the final forward method omits this fourth block and uses only the first three blocks.

After the convolutional feature extractor, we apply `AdaptiveAvgPool2d((1, 1))` to perform global average pooling, reducing each 256-channel feature map to a single scalar and producing a compact 256-dimensional representation per image. This global pooling approach reduces the number of parameters compared to flattening the entire spatial feature map and has a regularizing effect by enforcing spatial invariance. The classifier head then consists of a fully connected layer mapping from 256 units to 256 hidden units with a ReLU activation, followed by a dropout layer with probability 0.5, and a final linear layer mapping from 256 to C output units. The dropout layer is applied only in the classifier head, where overfitting is most likely due to dense connections and relatively few training examples per parameter.

We trained the model using the multi-class cross-entropy loss, which is appropriate for single-label classification tasks. For an input image with logits $x \in R^C$ and a one-hot label vector $y \in \{0, 1\}^C$, the loss is given by

$$L = - \sum_{i=1}^C y_i \log(\text{softmax}(x)_i),$$

where $\text{softmax}(x)_i$ denotes the predicted probability for class i . In practice, we use PyTorch’s `CrossEntropyLoss`, which combines the softmax and logarithm into a numerically stable implemen-

tation that directly consumes integer class indices. The optimization is performed using the Adam optimizer, initialized with a learning rate of 0.001 and no explicit weight decay in the final model. Adam adaptively scales the learning rate for each parameter based on first and second-order moment estimates of the gradients, which we found to converge faster and more reliably than stochastic gradient descent in our setting.

To further refine the optimization, we used a `ReduceLROnPlateau` learning rate scheduler in “max” mode, monitoring the validation accuracy. If the validation accuracy does not improve for three consecutive epochs, the scheduler reduces the learning rate by a factor of 0.5. This strategy allows the model to take relatively large steps in the early stages of training and gradually shift to smaller, more precise updates as the validation performance saturates. The training loop is structured around two functions: `train_step`, which runs one epoch over the training set (performing forward passes, computing loss and accuracy, and updating parameters), and `test_step`, which evaluates the model in eval mode on the test set without gradient updates. We track and store the training and test losses and accuracies for each epoch, and later visualize them using custom plotting functions.

We implemented and compared several model variants and baselines. As a simple non-neural baseline, a majority-vote classifier that always predicts the most frequent class in the training data would achieve an accuracy roughly equal to the relative frequency of that class (substantially below 50% on our relatively balanced dataset). Our CNN consistently outperforms this trivial baseline, with test accuracies in the 60–70% range for early configurations and approximately 84.75% for the best configuration. Within the neural models, we experimented with: (i) including the fourth convolutional block in the forward pass (increasing depth and channel capacity), (ii) adding additional dropout layers in the convolutional blocks, and (iii) introducing weight decay (L2 regularization) on the optimizer. While these ablations increased regularization and capacity, they often led to underfitting or reduced validation accuracy when combined with relatively strong data augmentation and limited dataset size. For example, using all four blocks together with dropout in each block and weight decay of 10^{-4} reduced the final validation accuracy to around 68%, indicating that the model became too constrained relative to the available data.

We also varied the input-side augmentation intensity, experimenting with heavier transformations such as color jitter and random rotations in addition to the random resized crop and horizontal flip. In some cases, these stronger augmentations modestly improved robustness but also slowed convergence and did not consistently exceed the performance of the simpler augmentation pipeline when paired with the three-block CNN. Based on these ablations, we selected the three-block architecture with a single dropout layer ($p = 0.5$) in the classifier, Adam optimization at an initial learning rate of 0.001, a ReduceLROnPlateau scheduler, and moderate spatial augmentation as our final model. This configuration strikes a balance between capacity and regularization, and it clearly outperforms the majority baseline and our earlier CNN variants in terms of validation accuracy on the municipal waste classification task.

5 Evaluation

The dataset was divided using an 80/20 split, with 80 percent being used for training and 20 percent being used for testing. This split was optimal because a higher percentage of testing data would result in the model being less accurate due to less data being used for training, and a lower percentage of testing data would result in overfitting because the testing dataset would not be representative of the data. The entire dataset was used instead of 30 percent to increase testing accuracy and reduce overfitting. The model was trained for 45 epochs using cross-entropy loss to determine the difference between the predicted class and the true label. We increased the number of epochs from 30 to 45 because increasing the number of epochs would increase the model’s accuracy, since the model will be able to accurately classify the images more often if it trains for a longer amount of time. We noticed that the model was not overfitting after 30 epochs, since the training accuracy and testing accuracy were similar, so training the model for a longer period of time would increase the training accuracy, which would lead to a higher testing accuracy. We found that the metrics from the progress report stage were adequate since they provided accurate results and were useful when tuning hyperparameters. The cross-entropy loss function used by PyTorch implements a softmax function for the predicted class and then multiplies the value by the true label for the class using the true label vector.

An Adam optimizer was used with a learning rate of 0.001, which is an algorithm that works similarly to stochastic gradient descent and updates the weights of the model using a dynamic learning rate. The algorithm starts with a learning rate of 0.001 but uses different learning rates as the algorithm determines what the ideal learning rate is. We did not change the learning rate because we noticed that it was high enough for the model to learn after 45 epochs, and the model was not overfitting, which meant that we did not need to increase the learning rate. Due to the size of the dataset and the simplicity of the model, the runtime to train and test the model was long.

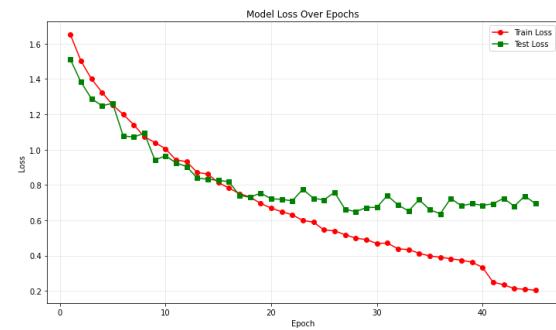


Figure 1: Model loss over 45 epochs for training and testing datasets

Using this methodology, the training loss for the first epoch was 1.65212, and the accuracy was 32.33 percent. This is a worse performance than our previous model for the first epoch, which had a training loss of 1.25839 and an accuracy of 52.17 percent. The testing loss for the first epoch was 1.51110, and the test accuracy was 41.60 percent, in comparison to the older model, which had a testing loss of 1.18205 for the first epoch and an accuracy of 54.86 percent. The data augmentation most likely led to the model being less accurate during earlier epochs. This demonstrated that the model in the first epoch performed poorly at determining the garbage type and that there was still a lot of room for improvement.

As we continued to train the model and as the number of epochs increased, the training and testing loss decreased, and the training and testing accuracy increased. This demonstrated that the model was improving at classifying garbage and was not overfitting to the dataset. This is because the increase in accuracy was steady, and the accuracy values were reasonable. By epoch 30, the training loss was 0.46969, which is less than the 0.69006 training loss in the older model, and the

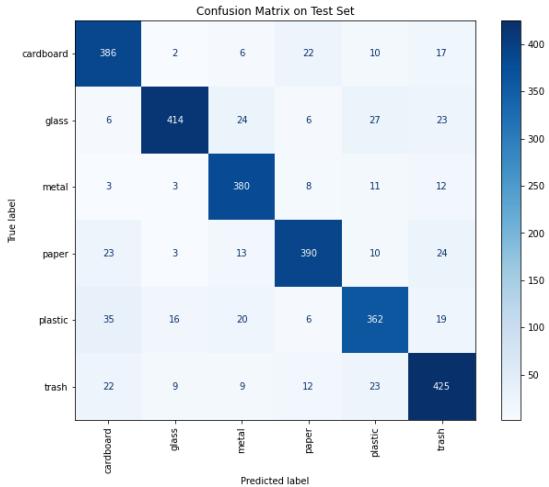


Figure 2: Confusion matrix for the model with 5 categories of waste

training accuracy was 83.29 percent, which was higher than the 75.65 percent training accuracy in the older model. Furthermore, the testing loss was 0.67411, which is less than the 0.92142 testing loss in the older model, and the testing accuracy was 80.05, which is greater than the 70.60 percent test accuracy in the older model.

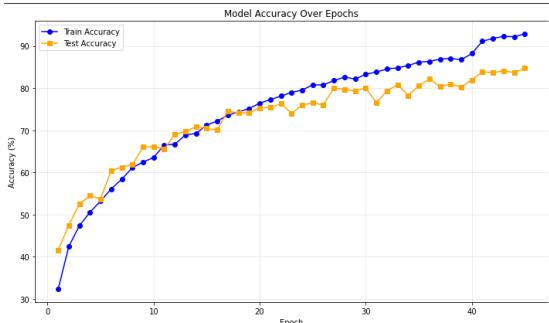


Figure 3: Accuracy over 45 epochs for training and testing datasets

This demonstrates that the older model was able to classify the waste fairly accurately, but the newer model was more accurate due to the changes to the dataset and the model. The training accuracy of the model after the 45th epoch was 92.88 percent, and the training loss was 0.20466. The testing accuracy was 84.75 percent with a testing loss of 0.69462. This meant that the model failed to classify images approximately 15 percent of the time, which was a huge improvement in comparison to the old model, which failed to classify images approximately 30 percent of the time. Similar to the old model, this model is also a great baseline for future models. Using a different architecture may minimally improve

the accuracy and reduce the runtime.

6 Progress

As we worked on this project, our model changed significantly as we refined our methodology based on the performance of the model. We initially looked into making large changes to our model, and looked into other models and how they perform with large datasets. However, we realized that the CNN is still a viable model to classify garbage, since it was an effective image processor due to its ability to learn relevant features easily on datasets with a low number of classes. The model is also able to work effectively with data augmentation, which is an important feature we needed to implement for our final model. We normalized our model on another model's pretrained mean and standard deviation since we noticed that our dataset was not normalized and could benefit from preprocessing.

We improved our preprocessing by implementing random crops and horizontal flips, which are commonly used data augmentation techniques. We used a ReduceLROnPlateau learning rate scheduler to adjust the learning rate when the model accuracy does not improve, which improved the testing accuracy. This demonstrated that we followed through with our plan to adjust our model. However, our plan did change course when we also experimented with adding a 4th convolutional block to make the network deeper, adding additional dropout layers, and including L2 regularization into our convolutional neural network. This is because we ended up not implementing these features as planned, since they did not improve the model's accuracy and led to underfitting.

7 Error Analysis

A thorough review of the model's mistakes involved quantitative tools such as confusion matrices and per class accuracy, along with qualitative checks that included visualizing intermediate CNN layers to see which features the network relied on when sorting municipal waste into six categories. These visualizations helped us identify when the model focused too much on background patterns or irrelevant textures, which signaled overfitting and guided how we tuned dropout and other hyperparameters. The model generally performed well on categories with strong visual cues, such as metal or glass, while struggling with classes that looked similar in color or texture, such as certain plastics and

paper items. Comparing different model versions showed clear shifts in what errors decreased after architectural changes or regularization adjustments. Common error patterns included confusion among items with similar shapes and high confidence on ambiguous samples. Future improvements could involve more targeted data collection, stronger augmentation, class balancing, or model components that better highlight material specific features.

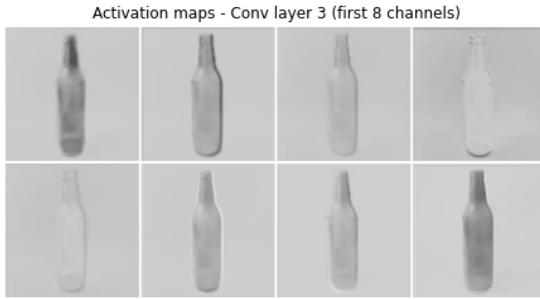


Figure 4: Example convolutional layer activation visualization for glass category

We also used a McMaster server to run parallel grid searches over batches of hyperparameters to speed up experimentation and improve both accuracy and convergence. Running these searches in parallel let us evaluate many combinations of learning rates, dropout values, batch sizes, and optimizer settings without long wait times. This setup made it easier to spot trends in what configurations produced faster training, more stable losses, or better validation accuracy. It also helped confirm when certain tuning choices, such as increased regularization or adjusted learning rate schedules, consistently reduced overfitting across multiple runs.

Team Contributions

Yousef Shahin: Completed section 3, and 4 (Features and Inputs, and Implementation), created the model architecture and training pipeline, and participated in trainig the model and performing experiments.

Aswin Kuganesan: Completed Dataset, Evaluation, and Progress sections. Created data augmentation and tested different model architectures.

Sunny Yao: Completed sections Introduction, Progress and Error Analysis. Created the confusion matrix generation and randomized sampling and CNN feature visualizations.